

02-05: Vector Operations and NA

1 - Purpose

- mathematical operations on vectors
- dealing with missing values in a vector
- R shortcuts for vector mathematics

2 - Questions about the material...

If you have any questions about the material in this lesson or your Class Project *feel free to email them to the instructor [here](#)*.

3 - Vector math

For this lesson, we are going to use the data from the file **LansingWeather2.csv**. First we need to open the CSV file and save the data in the file to a data frame, which we will call ***weatherData***.

```
1 {  
2   rm(list=ls()); options(show.error.locations = TRUE);  
3  
4   # read data from LansingWeather2.csv and save to the variable weatherData...  
5   weatherData = read.csv("data/LansingWeather2.csv");  
6 }
```

The data frame, ***weatherData***, has the date, high temperature, low temperature, and precipitation for 14 days (March 27-April 9, 2017) in Lansing, MI.

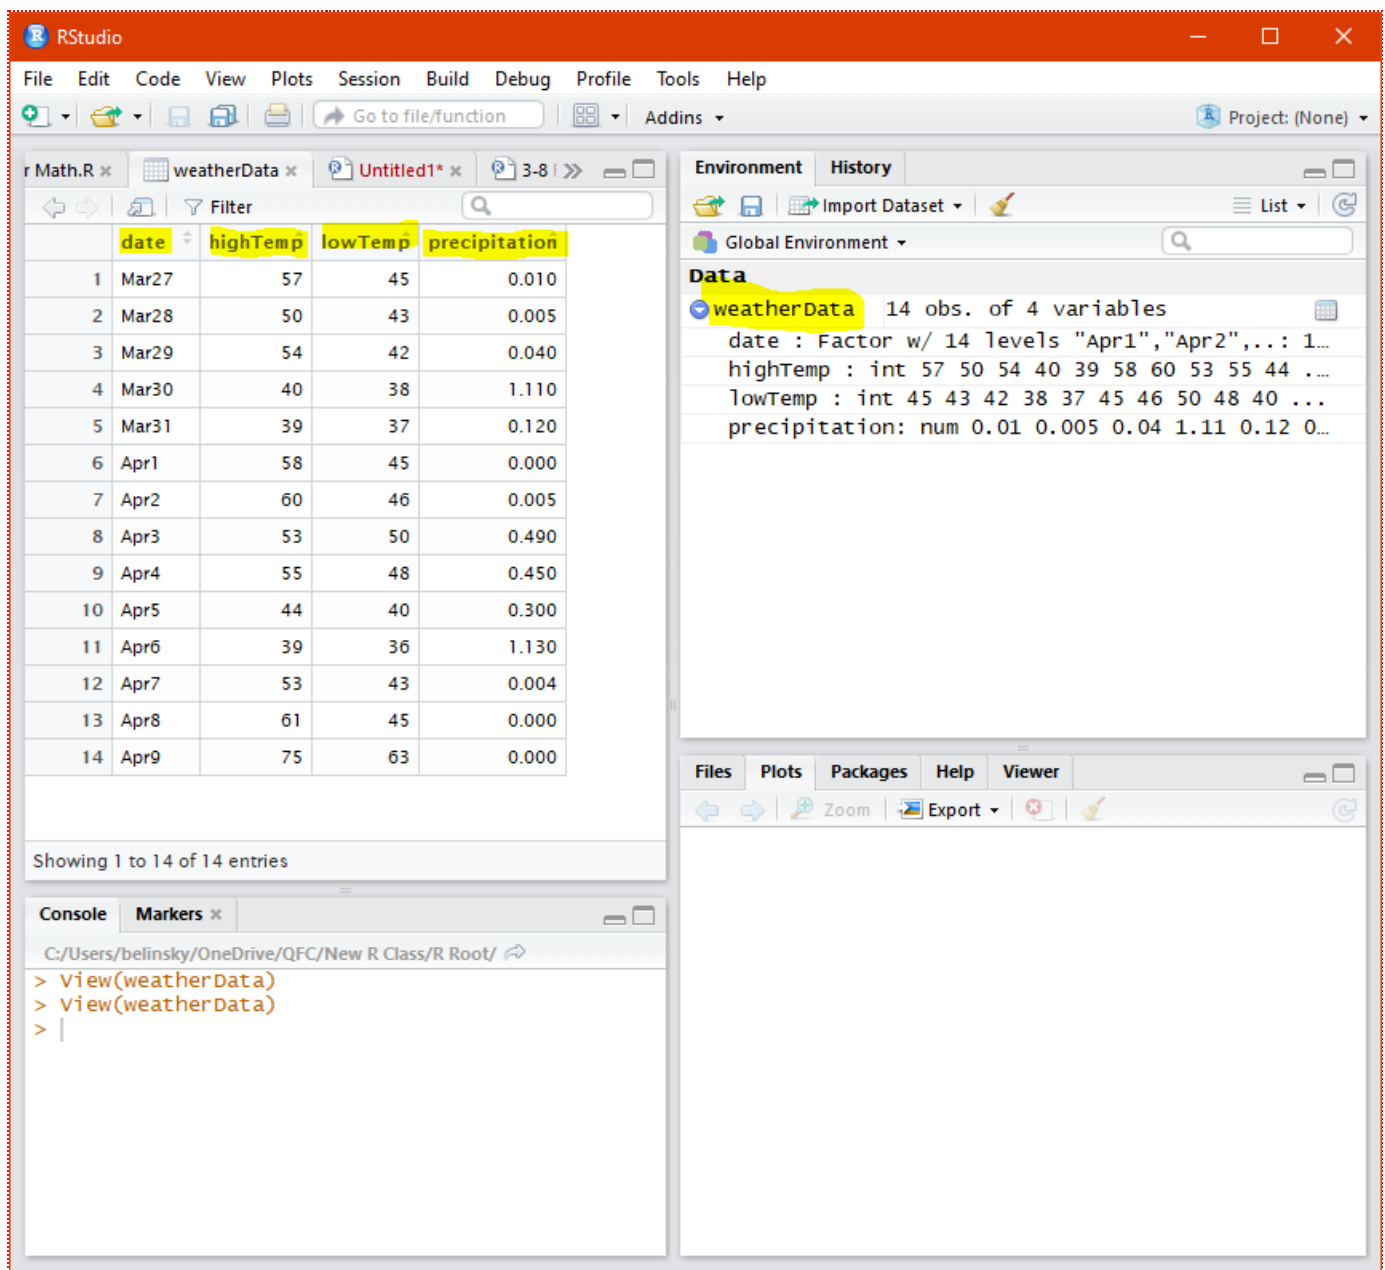


Fig 1: Looking at the values in the **weatherData** data frame in the Main Window.

3.1 - Finding change in temperature

We want to find the temperature range for each day -- in other words, the high temperature minus the low temperature. This means subtracting each day's low temperature from the high temperature and saving that value to a vector.

First we will extract the low and high temperature values from **weatherData** and save them to a vector. There are two ways we can do this:

1) by column number

```
1 # get all values from the 3rd column in weatherData
2 highTemp = weatherData[, 3];
```

2) or by column name

```
1 # get all values from the column named "lowTemp" in weatherData
2 lowTemp = weatherData[, "lowTemp"];
```

Note: the row number is left blank to indicate we are getting values *from all rows*.

We will use a **for()** to iterate through each value (i.e., the highTemp - lowTemp for each day) and solve for the temperature difference. To do this we also need a vector that will hold the change in temperature values, which we will call **changeInTemp**.

```
1 changeInTemp = c(); #declared a vector
```

changeInTemp acts as a state variable because:

- 1) **changeInTemp** gets initialized before the **for()**
- 2) **changeInTemp** gets populated during the iterations of the **for()**
- 3) **changeInTemp** final state is the full set of temperature differences

3.2 - Subtracting values from two different vectors

We are going to use a **for()** loop that iterates through each value in **highTemp** and **lowTemp**, finds the difference, and saves the answer to **changeInTemp**. We need to know the length of the vectors to do this. In this case, we already know there are **14** values, but we will use **length()** to calculate this value so that it works for vectors of any size (i.e., if we wanted to run the same code but use more days).

```
1 vectorLength = length(lowTemp);
```

Note: since the length of all columns in a data frame are by definition the same, we only need to get the length of one vector.

We use **vectorLength** to create a sequence that the **for()** iterates through. The **for()** iterates through the sequence **1:vectorLength**, (**1:14**), and assigns the value of the **14 highTemp - lowTemp** difference operations to the **14 changeInTemp** vector elements.

```
1 for(i in 1:vectorLength)
2 {
3   changeInTemp[i] = highTemp[i] - lowTemp[i];
4 }
```

Let's put all this code together:

```
1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   # read data from Lansingweather2.csv and save to variable weatherData
5   weatherData = read.csv("data/Lansingweather2.csv");
6   # get all values from the 3rd column in weatherData
```

```

7 highTemp = weatherData[, 2];
8 # get all values from the column named "lowTemp" in weatherData
9 lowTemp = weatherData[, "lowTemp"];
10
11 changeInTemp = c(); # declare a vector
12 vectorLength = length(lowTemp); # vectorLength will be 14 (length of data)
13
14 # go through the sequence 1:14
15 for(i in 1:vectorLength)
16 {
17     # subtract lowTemp from highTemp for all 14 values and save to changeInTemp
18     changeInTemp[i] = highTemp[i] - lowTemp[i];
19 }
20 }

```

The Environment Window shows the vector **changeInTemp** but it only shows up to 10 values. If you want to see all the values in **changeInTemp**, you can type **changeInTemp** in the Console Window and the 14 values in **changeInTemp** will appear on the next line.

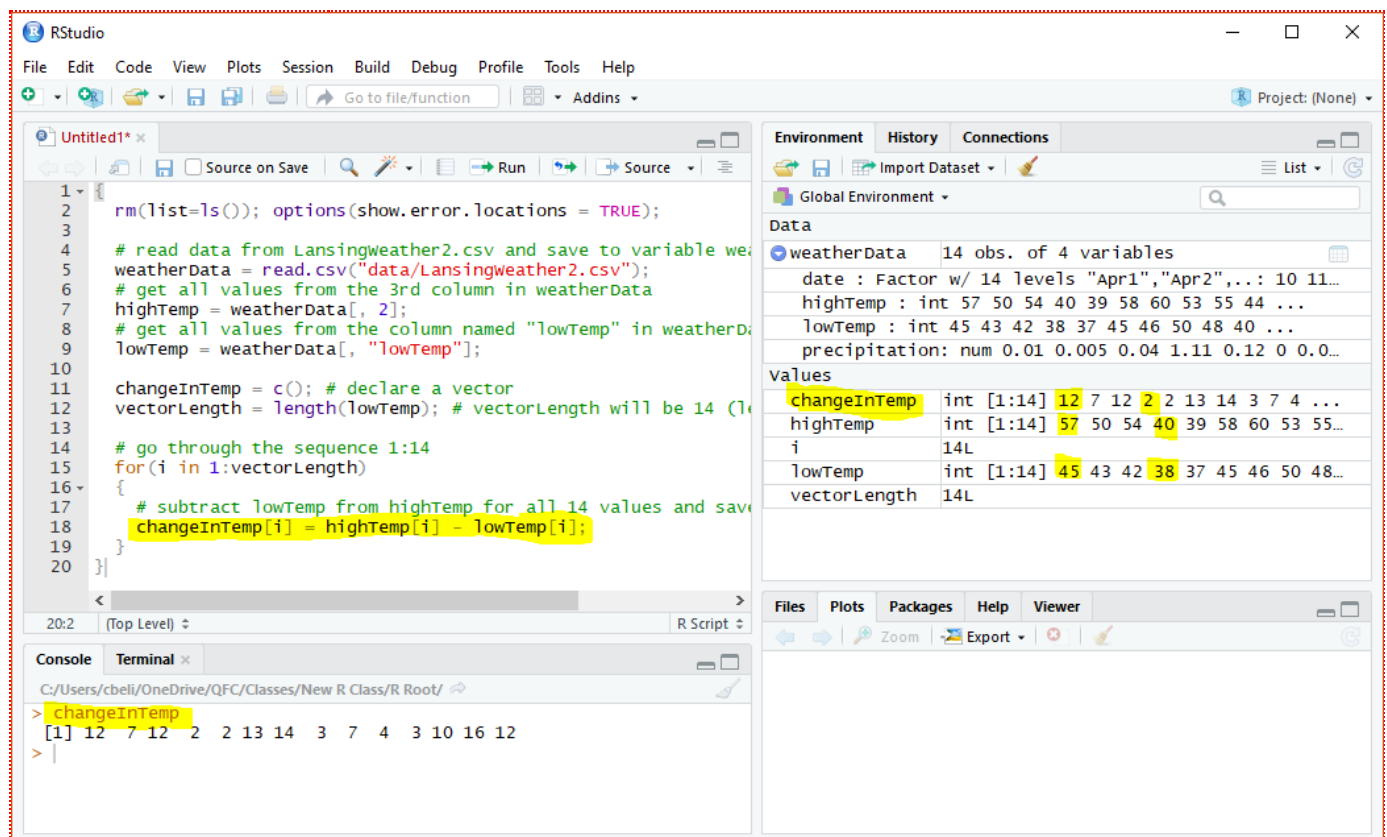


Fig 2: Using a **for()** to perform iterative mathematical operations on vectors.

4 - Dealing with missing values

In the previous example, we have an idealized situation where every day had a high and low temperature associated with it. In the real world, especially with large amounts of data, there are often missing data. In R,

missing data are represented by **NA**, but the CSV file could designate missing data according to a number of different conventions.

First lets create a data set with missing values and save this as **MissingTemps.csv** in the **Data** directory

```
1 date,highTemp, lowTemp, precipitation
2 Mar27,57,45,0.01
3 Mar28,50,43,0.005
4 Mar29,54, ,0.04
5 Mar30,40,38,1.11
6 Mar31,39,NA,0.12
7 Apr1,58,45,0
8 Apr2,60, ,0.005
9 Apr3,53,50,0.49
10 Apr4,55,48,0.45
11 Apr5,44,40,0.30
12 Apr6,39,36,1.13
13 Apr7,NULL,43,0.004
14 Apr8,61,45,0
15 Apr9,75,63,0
```

In the above data set, the **lowTemp** for **Mar29** and **Apr2** are left blank, the **lowTemp** for **Mar31** is given as **NA**, and the **highTemp** for **Apr7** is given as **NULL**. So, there are four missing values in **missingTemps.csv**.

4.1 - Viewing blank or NA values

We are going to open a new script and save **missingTemps.csv** data to a Data Frame called **weatherData**, but we want to standardize the way that missing values are recorded. So, in **read.csv()** we add the parameter **na.strings**:

```
1 na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null")
```

Essentially, **na.strings** is a vector that contains all the value that you want R to assign as **NA**. It is good to be paranoid here and think of all the possible ways in which CSV files will present these values! In this case, we only needed to use "", "NULL", and "NA".

```
1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   weatherData = read.csv("data/missingTemps.csv",
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));
6 }
```

Execute the code above and double-click on **weatherData** in the **Environment Window** so that it appears in the Main Window. We see that all four missing values in **weatherData** are now labelled as **NA**.

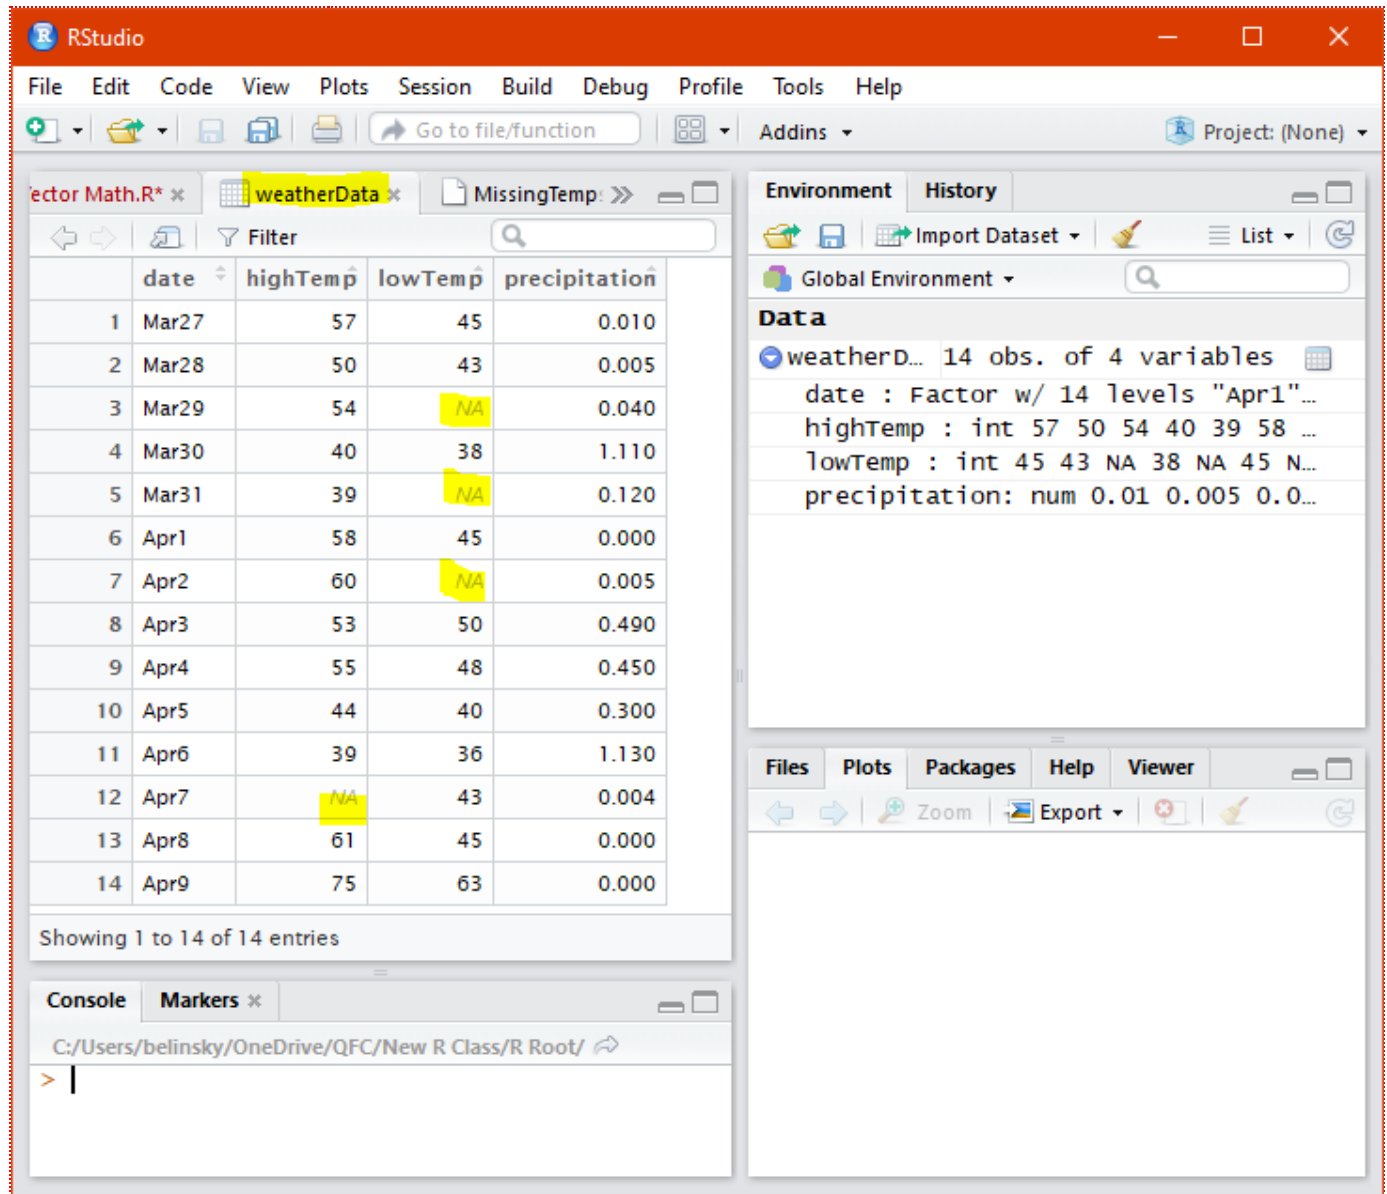


Fig 3: NA values in a data frame

4.2 - Mathematical operations on vectors with NA values

We are going to run the same change in temperature script as before except with NA values in the data frame:

```
1 {  
2   rm(list=ls()); options(show.error.locations = TRUE);  
3  
4   # read data from missingTemps.csv and save to the variable weatherData...  
5   weatherData = read.csv("data/missingTemps.csv",  
6     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));  
7
```

```

8 # get all values from the 3rd column in weatherData
9 highTemp = weatherData[, 2];
10 # get all values from the column named "lowTemp" in weatherData
11 lowTemp = weatherData[, "lowTemp"];
12
13 changeInTemp = c(); # declare a vector
14 vectorLength = length(lowTemp); # vectorLength will be 14 (length of data)
15
16 # go through the sequence 1:14
17 for(i in 1:vectorLength)
18 {
19     # subtract lowTemp from highTemp for all 14 values and save to changeInTemp
20     changeInTemp[i] = highTemp[i] - lowTemp[i];
21 }
22 }

```

Note: anytime there is an **NA** in a calculation (e.g., 3rd element in **changeInTemp**), the answer is going to be **NA**. This is a special feature of **NA** -- if you used any other value to represent missing data then you would get an error in the calculation.

The screenshot shows the RStudio interface. The script editor on the left contains the code from the previous block. The Environment pane on the right shows the following data:

Variable	Class	Length	Values
weatherData	data.frame	14 obs. of 4 variables	
date	Factor	w/ 14 levels	"Apr1", "Apr2", ...: 10 11 ...
highTemp	int	[1:14]	57 50 54 40 39 58 60 53 55 44 ...
lowTemp	int	[1:14]	45 43 NA 38 NA 45 NA 50 48 40 ...
precipitation	num		0.01 0.005 0.04 1.11 0.12 0 0.00...
changeInTemp	int	[1:14]	12 7 NA 2 NA 13 NA 3 7 4 ...
highTemp	int	[1:14]	57 50 54 40 39 58 60 53 55...
i	int	14L	
lowTemp	int	[1:14]	45 43 NA 38 NA 45 NA 50 48...
vectorLength	int	14L	

The Console pane at the bottom shows the output of the code execution:

```

> source("~/active-rstudio-document")
> changeInTemp
[1] 12 7 NA 2 NA 13 NA 3 7 4 3 NA 16 12

```

Fig 4: Doing mathematical operation with **NA** values

Extension: What counts as **NA**

5 - Vector Operation Shortcuts

So far we have used **for()** loops to perform iterative operations on a vector or multiple vectors. However, R has built in functions that do all these operations using a lot less code. For example, we can perform the above change in temperature calculation without a **for()**.

```
1 {  
2   rm(list=ls()); options(show.error.locations = TRUE);  
3  
4   weatherData = read.csv("data/missingTemps.csv",  
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"));  
6  
7   highTemp = weatherData[, 2];  
8   lowTemp = weatherData[, "lowTemp"];  
9  
10  changeInTemp = highTemp - lowTemp;  
11 }
```

Line 9 does all the work of iterating through the values in the vector, subtracting the values, and saving the answer to the vector **changeInTemp**. The result is the same as the script that uses the **for()** loop.

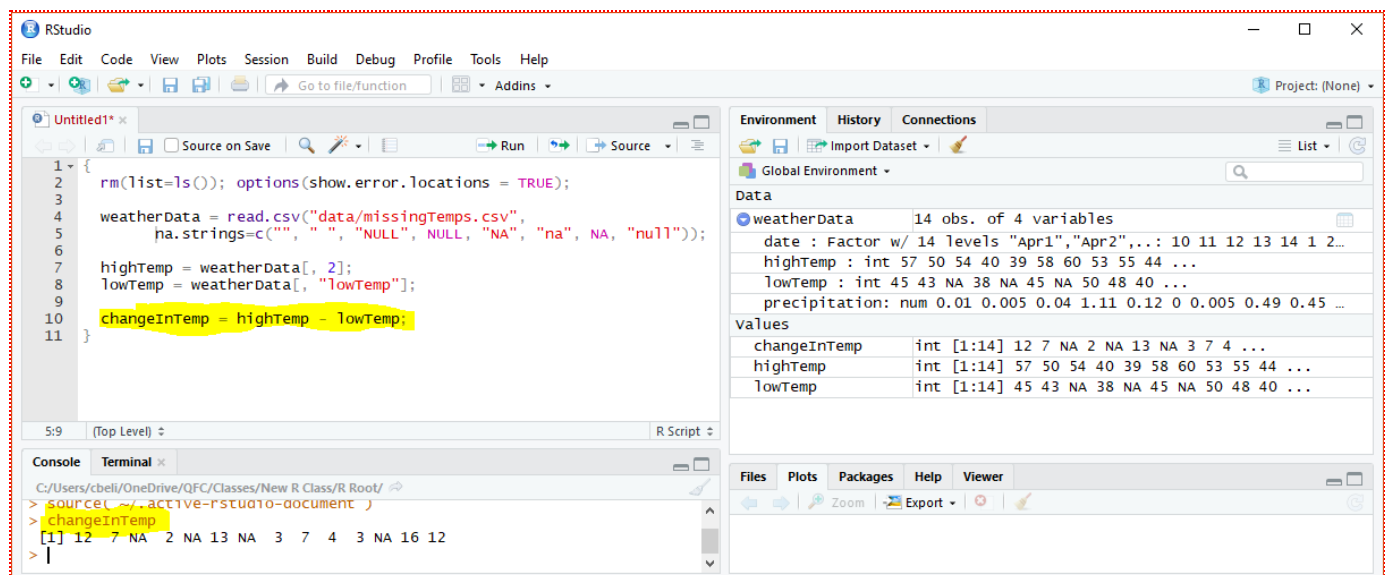


Fig 5: Subtracting a vector in R.

This method is obviously easier than **for()** loops for subtracting two vectors -- and it is the method R programmer will usually use. However, this method is really a shortcut and it is important to understand what is going on in with the **for()** example because, eventually, you will find a situation where you need to use a **for()** to solve a more complex problem.

5.1 - Many other shortcuts

R has many functions that can quickly perform the most common operations on vectors like **sum()**, **max()**, **min()**, and **mean()**.

One important parameter used in all of these functions is **na.rm**:

na.rm = TRUE tells R to exclude the **NA** values from the vector before performing the operation.
na.rm = FALSE tells R to return **NA** if there are **NA** values in the vector.

The default value for ***na.rm*** is ***FALSE***.

```
1 {
2   rm(list=ls()); options(show.error.locations = TRUE);
3
4   weatherData = read.csv("data/missingTemps.csv",
5     na.strings=c("", " ", "NULL", NULL, "NA", "na", NA, "null"))
6
7   # save high temp and low temp columns to vectors
8   highTemp = weatherData[, 2];
9   lowTemp = weatherData[, "lowTemp"];
10
11  # get min and max temp (na.rm = default = FALSE)
12  minTemp = min(lowTemp);
13  maxTemp = max(highTemp);
14
15  # get min and max temp ignoring the NAs
16  minTempTake2 = min(lowTemp, na.rm=TRUE);
17  maxTempTake2 = max(highTemp, na.rm=TRUE);
18
19  # get the index of min and max values
20  minIndex = which.min(lowTemp);
21  maxIndex = which.max(highTemp);
22
23  # get simple statistics on a vector
24  sumTemp = sum(highTemp, na.rm=TRUE);
25  meanTemp = mean(highTemp, na.rm=TRUE);
26  medianTemp = median(highTemp, na.rm=TRUE);
27  stanDev = sd(highTemp, na.rm=TRUE);
28  variance = var(highTemp, na.rm=TRUE);
29
30  # get index of values that meet a condition
31  # this is more advanced and will be covered in detail later
32  whichHighGT60 = which(highTemp > 60);
33  whichLowLT43 = which(lowTemp < 43);
34 }
```

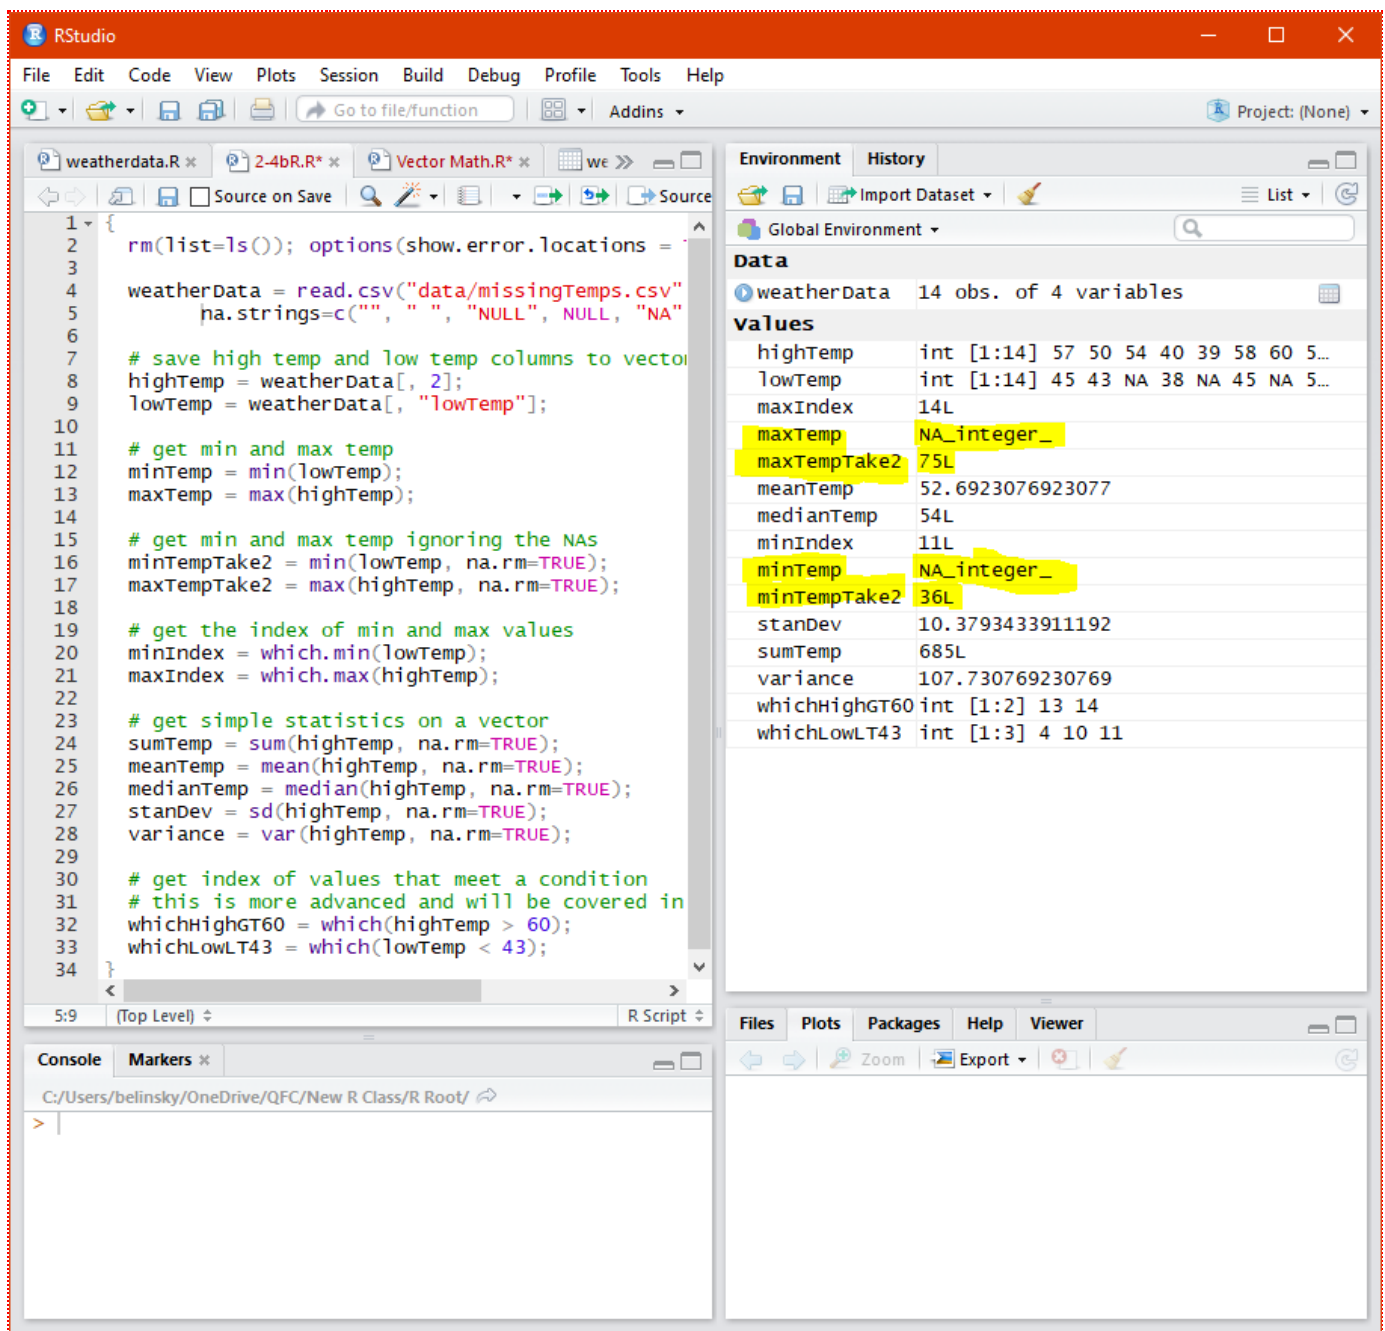


Fig 6: Using R functions to perform common vector operations.

6 - Application

If you have any questions regarding this application or your Class Project, feel free to email them to the instructor here. You can attach the whole Root Folder as a zipped file.

1) Create a vector called **changeInHighTemp** and, using the data from **lansingWeather2.csv**, find the change in high temperatures from day-to-day.

So, if the four high temperatures are: **40, 45, 35, 42**

then the three **changeInHighTemp** would be: (45-40), (35-45), and (42-35) = **5, -10, 7**

2) There will be one less value in **changeInHighTemp** than in **lansingWeather2.csv**. We want the length of **changeInHighTemp** to be the same so we can add it to the data frame (remember, all vectors must be the

same size in a data frame). To do this, add an **NA** value as the first value in **changeInHighTemp**. The **NA** indicates that there was no previous value to subtract and, hence, this value is not applicable.

3) Add the vector **changeInHighTemp** (with the **NA** as the first value) to the data frame.

*Save your script file as **app2-5.r** in the **scripts** folder of your RStudio Project for the class.*

7 - Extension: What counts as NA

R is somewhat savvy when it moves data from a CSV to a data frame. If there is an empty value in a column that is all numbers, then R will convert the empty value to **NA**. **NULL**, however, is not considered to be a valid data frame value and if there is a **NULL** in a column, it will be treated as the string value "NULL". R also does not consider lowercase **na** to be **NA** and this will also be treated as the string value "na".

There is now a second problem: if there is even one string in a column, all values in the column are treated as strings. This means that **40** will be seen as the string "40" and mathematical operations on strings will cause an error. Thus, specifying the **na.strings** argument can be useful to ensure that you will not have to convert columns to numeric using an extra step after the data are imported.